


Multimodal RAG Reference Architecture

This whitepaper describes a reference architecture for retrieval-augmented generation over multimodal document collections. The pipeline is organized into five decoupled stages: ingestion, embedding, indexing, retrieval, and generation, as illustrated in the architecture diagram on the following page.

Stage	Primary Technology
Ingestion	PyMuPDF, pdfplumber, pytesseract
Embedding	CLIP (sentence-transformers)
Vector Store	ChromaDB
Generation	GPT-4V-class VLM

Figure: Supporting Chart

Data Flow: Multimodal RAG Pipeline



The chart above visualizes the key figures referenced on page 1.

Figure: Additional Diagram

Data Flow: Multimodal RAG Pipeline

